


Decoding the Vulnerability Backlog

How Explainable NLP and Predictive Synthesis are Restructuring the Software Supply Chain

TARGET AUDIENCE: Cybersecurity Leadership // Data Scientists // Engineering Leads



The NVD Choke Point

The Reality


The NVD is buckling under the exponential surge of software supply chain flaws.

Vulnerabilities are arriving faster than NIST can manually analyze and enrich them.

The Stance

NIST acknowledges this is "very labor-intensive" and "not scalable."

They are currently "fighting a losing battle."




Exponential Surge of Supply Chain Flaws


The Impact

Incomplete vulnerability reports severely hinder automated threat mitigation, compliance auditing, and patching workflows.

The Policy Shift: Distributing the Bottleneck



Centralized Manual Enrichment



Decentralized Enrichment Pipeline

The Policy Shift

NIST intends to transfer vulnerability-enrichment work to CNAs and prioritize triage (focusing on the CISA KEV catalog and critical software).

The Ecosystem Response

CISA has launched "Vulnrichment" and Europe launched the Global CVE Allocation System (GCVE) to compensate, risking a balkanized process.


The Core Tension

Decentralizing manual analysis to resource-constrained CNAs simply distributes the bottleneck; it does not eliminate it. A "large reset" requires a new technological approach.


The Anatomy of the Void: Statistical Bias in Registries

Unenriched CVEs don't just create operational lag; they create statistically biased datasets. Understanding why data is missing dictates how we build predictive mitigation models.


The Typology of Missingness



MCAR



MAR



MNAR

Missing Data Diagnostic Matrix


	Cybersecurity Context	Downstream Impact	ML Resolution Strategy
MCAR	API timeouts, network transmission glitches, or administrative database corruption.	Reduces statistical power but does not introduce parameter bias.	Complete-case analysis (listwise deletion) or pairwise deletion.
MAR	Resource-constrained CNAs systematically failing to supply CVSS vectors or CWE classifications.	Introduces severe estimation bias if ignored; standard statistics fail.	Multiple imputation, likelihood-based factoring, or deep learning classifiers.
MNAR	Proprietary vendors withholding exploit details; low-severity flaws under-reported by maintainers.	Severe underestimation of risk; systemic bias in severity scoring.	Advanced cost-sensitive classifiers, domain-specific text models (NLP).

The Foundation of Truth: Dataset Replicability

The Readiness Crisis

Machine learning requires replicable, high-quality datasets. The VALIDATE study systematically reviewed 50 primary vulnerability prediction studies to assess the viability of their ground truth.


The community is frequently reinventing the wheel with unstable, ad-hoc datasets, hindering the development of reliable predictive models.




Dataset Viability: The Pivot to Text Mining

Current Dataset Characteristics


Granularity of Labeled Entity



Nature of Datasets



Type of Labelling



NLP > Software Metrics

A 100-times 3-fold cross-validation replica study confirmed that **Text/String features (NLP) deliver significantly better performance and inspection ratios** than traditional software metrics.

(p-value < 0.0001)

The Mislabeling Threat

71% of gathered datasets do not assess negative labels.

71%


The Single-CWE Fallacy

Of the 29% that DO assess negatives, they only verify the absence of ONE specific CWE. A fragment deemed “negative” only means it lacks CWE-321; it does not mean it is free of all vulnerabilities.

The Impact

Consequently, past ML models and human analyses have systematically overestimated the false positive rate.

The Solution: Predictive Synthesis via NLP



Predictive Synthesis

We leverage NLP to automatically determine the CVSS vector directly from the textual description at disclosure.


VulnerWatch Efficacy

Deploying eight fine-tuned BERT models achieves favorable accuracy across metrics, averaging close to 0.9.

Severity Score Performance

Mean Squared Error (MSE) = 1.27 | Mean Absolute Error (MAE) = 0.5 | R2 = 0.51

Explainable NLP Architecture



Opening the Black Box

Gradient-based input saliency methods isolate the exact input words driving the model's prediction.


Aligning with Human Rationale

The top relevant words highlighted by the algorithm consistently map directly to the technical rationales used by human cybersecurity experts, ensuring the output is completely comprehensible.

Manual vs. Predictive Enrichment Matrix


	Manual Analysis (Legacy)	Predictive Synthesis (CVSS-BERT)
Speed to Triage	Days to Weeks	Near Real-Time at Disclosure
Scalability	Linear (Failing under volume)	Exponential (Automated via API)
Bias Handling	Highly susceptible to MNAR/MAR omissions	Algorithmic imputation based on global text patterns
Explainability	Dependent on inconsistent human notes	Consistent, gradient-based word saliency mapping

Future State Ecosystem Workflow



NIST cannot successfully shift enrichment responsibility to CNAs unless those CNAs are armed with automated NLP triage tools. CVSS-BERT is the critical enabler of a decentralized vulnerability ecosystem.

**The NVD backlog is not a failure of analysis;
it is a failure of scale.**



**By arming the software supply chain with explainable NLP,
we transform unstructured chaos into structured, predictive defense.**